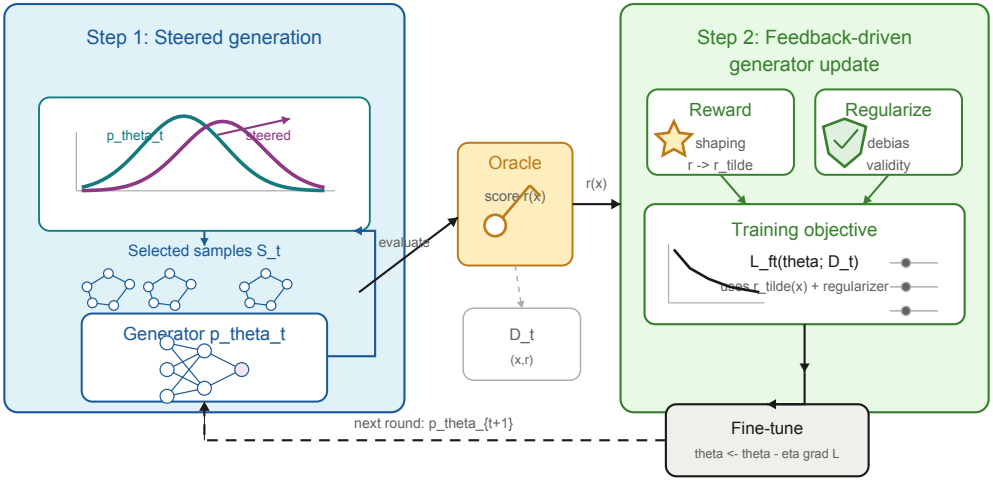
Step 1: Steered generation
p_theta_t
steered
Selected samples S_t
Generator p_theta_t
evaluate
Oracle
score r(x)
r(x)
D_t
(x,r)
Step 2: Feedback-driven generator update
Reward
shaping
r -> r_tilde
Regularize
debias
validity
Training objective
L_ft(theta; D_t)
uses r_tilde(x) + regularizer
Fine-tune
theta <- theta - eta grad L
next round: p_theta_{t+1}